

NoSQL Assignment 3

IMT2022001 Siddharth Menon
IMT2022006 Shreyas Arun Saggere
IMT2022082 Sathvik S. Rao
IMT2022507 Vishruth Vijay

April 14, 2025

Question 1:

The objective in this question was to define a Hive schema for the data files given, and create Hive tables accordingly. We initially created a staging table in order to load the data as it is. Then we loaded the data following our schema. We also defined an error schema for erroneous records. Our Hive schema looks as follows:

1. **Course_Attendance.csv** The **course_attendance_staging** table was used for staging.

Column Name	Data Type	Description
course	TEXT	Course identifier or name
name_hash	TEXT	Hashed student name
email_hash	TEXT	Hashed email address
member_id_hash	TEXT	Hashed member ID
classes_attended	INTEGER	Number of classes attended
classes_absent	INTEGER	Number of classes missed
instructors	ARRAY<String>	List of instructors (parsed from string)
avg_attendance_percent	DECIMAL(5,2)	Average attendance percentage

Table 1: Schema of the **course_attendance** Table

2. **Enrollment_Data.csv** The `enrollment_data_staging` table was used for staging.

Column Name	Data Type	Description
serial_no	INTEGER	Serial number
course	TEXT	Course name
status	TEXT	Enrollment status
course_type	TEXT	Type of course
course_variant	TEXT	Course variant
academia_lms	TEXT	NA
student_id	TEXT	Student ID
student_name	TEXT	Name of the student
program	TEXT	Program enrolled in
batch	TEXT	Batch or cohort
period	TEXT	Academic period or semester
enrollment_date	DATE	Date of enrollment
primary_faculty	ARRAY<TEXT>	List of faculty members for the course

Table 2: Schema of the `enrollment_data` Table

3. **grade_roster_report.csv**

Column Name	Data Type	Description
academy_location	STRING	Location of the academy or campus
student_id	STRING	Unique identifier for the student
student_status	STRING	Status of the student
admission_id	STRING	Admission ID associated with the student
admission_status	STRING	Status of the admission
student_name	STRING	Full name of the student
program_code_name	STRING	Program code and name
batch	STRING	Batch or cohort information
period	STRING	Academic period or term
subject_code_name	STRING	Subject code and name
course_type	STRING	Type of course
section	STRING	Section identifier
faculty_name	STRING	Name of the faculty/instructor
course_credit	INT	Credit value of the course
obtained_marks_grade	STRING	Marks or grade obtained by the student
out_of_marks_grade	STRING	Maximum marks or grade possible
exam_result	STRING	Exam result

Table 3: Schema of the `grade_roster_report` Table

4. Erroneous record schema

Column Name	Data Type	Description
source_table	STRING	Name of the source table where the error originated
original_row	STRING	Raw data row containing the error
column_name	STRING	Name of the column that caused the error
issue_type	STRING	Type or category of the issue (e.g., NULL, format error)
error_description	STRING	Detailed description of the error

Table 4: Schema of the `error_records` Table

Question 2

The `course_attendance table`, `enrollment_data table`, `grade_roster_report` table were joined on the fields `student_id` and `course` field to create a warehouse filtering out erroneous data. This warehouse was used in order to perform our analytical queries.

The analytical queries were:

1. **Query 1:** Comparing the mode grade and average attendance of different students of different programs for every course
2. **Query 2:** Percentage of students with a below 50% attendance record, per program, per course.
3. **Query 3:** Identification of top 3 courses by average attendance for each program.

Question 3

We set up a new data warehouse that is partitioned by `batch/program` and bucketed into 4 buckets by `course`. The goal was to evaluate the impact of partitioning on query performance. We chose to partition by `program` because our queries frequently include a ‘GROUP BY’ operation on the program field. For the same reason we chose to bucket by `course`.

Question 4: Converting Hive queries to Pig

Important note: A difference in schema between Hive and Pig: in the Hive schema, we have two fields `batch` and `period` – which would contain values like (Master of Technology, 2024–2026 CSE), while these same fields are renamed to `program_name` and `batch`, where `batch` is analogous to the `period` field in the Hive schema. For this section, those two fields will be referred to by the names they are referred to in the Pig schema.

Furthermore, since the code for each of the queries was put in a different file and executed one after the other from the terminal, the load had to be done in all the three query files, causing an additional overhead for a load every time. If the 3 queries are executed together, and not timed separately, the time taken is lesser, as we load only once. That analysis, too, has been done.

1. Comparing the mode grade and average attendance of different students of different programs for every course.

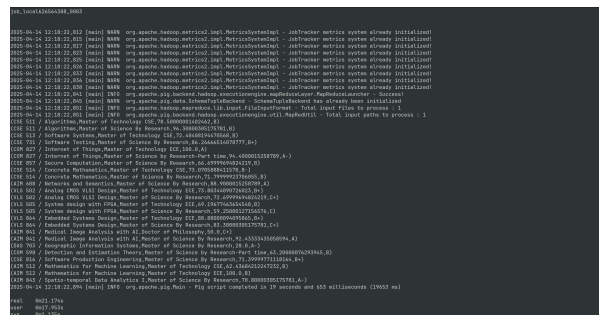


Figure 2: Time taken for query 1


[illegible]

Figure 4: Time taken for query 2

4

[illegible]

Figure 6: Time taken for query 3

Executing all the 3 queries together and then timing it.

[illegible]

Figure 8: All queries together - time taken

Query	Execution Time (seconds)
Query 1	21.174
Query 2	22.108
Query 3	15.485
All Queries Executed Together	50.720

Table 5: Execution times of individual queries and combined execution

1 Comparing the performance of Hive(Partitioned and non-partitioned) and Pig

Query 1: Average attendance and mode grade per course per batch

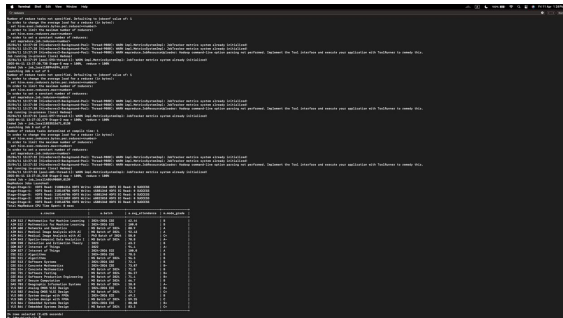


Figure 9: Query 1 - Non-Partitioned Execution

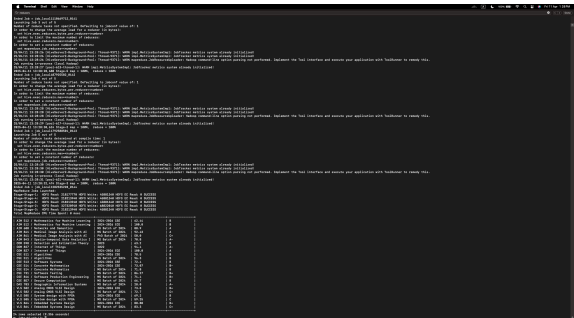


Figure 10: Query 1 - Partitioned Execution

Query 2: Percentage of students with below 50% attendance per course and program

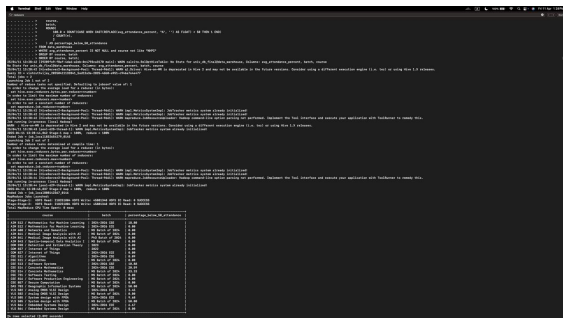


Figure 11: Query 2 - Non-Partitioned Execution

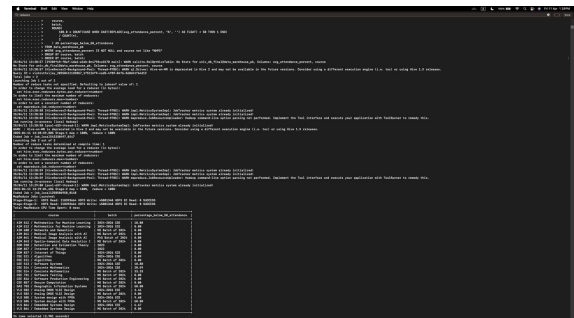


Figure 12: Query 2 - Partitioned Execution

Query 3: Top 3 courses by average attendance per program

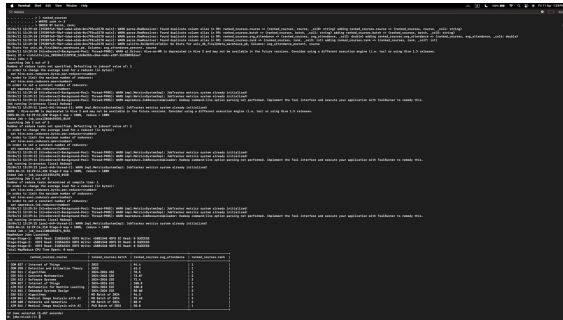


Figure 13: Query 3 - Non-Partitioned Execution

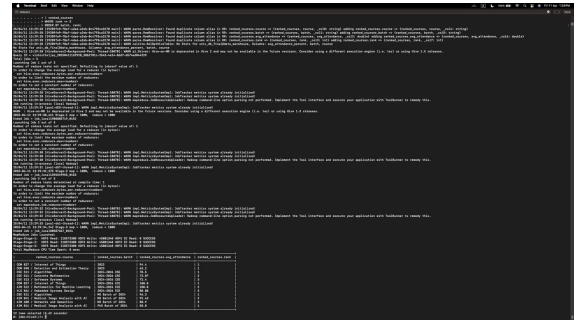


Figure 14: Query 3 - Partitioned Execution

Execution Time Comparison Table

Query	Partitioned + Bucketed Time (s)	Non-Partitioned Time (s)	Pig (s)
Query 1	9.356	9.62	21.174
Query 2	3.892	3.901	22.108
Query 3	5.62	5.657	15.485

Table 6: Execution Time Comparison: Partitioned vs Non-Partitioned Data Warehouse